

Module 05: Model Evaluation

Learning Objectives:

- Model Evaluation
- Over-fitting, Under-fitting and Model Selection
- Ridge Regression
- Grid Search
- Question:

* How can you be certain your model works in real world and performs optimally.

Model Evaluation

- In-sample evaluation tells us how well our model will fit the data used to train it?
- Problem?
 - It does not tell how well the trained model can be used to predict new data
- Solution?
 - In-sample data or training data,
 - Out-of-sample evaluation or test set

$y_{\text{data}} \leftarrow$ dataset target; $\text{df}[\text{price}]$
 $x_{\text{data}} \leftarrow$ features or independent variables

$X_{\text{train}}, x_{\text{test}}, y_{\text{train}}, y_{\text{test}} = \text{train-test-split}(x_{\text{data}}, y_{\text{data}}, \text{size} = 0.3)$
random-slice = 0

from sklearn.model_selection import train_test_split

• Split data into random train and test subsets

Under train-test-split()

the idea is to train the model to get the best performance.
When we have completed testing our model we should use all
use testing set to assess the performance of a predictive model.
Build and train the model with a training set

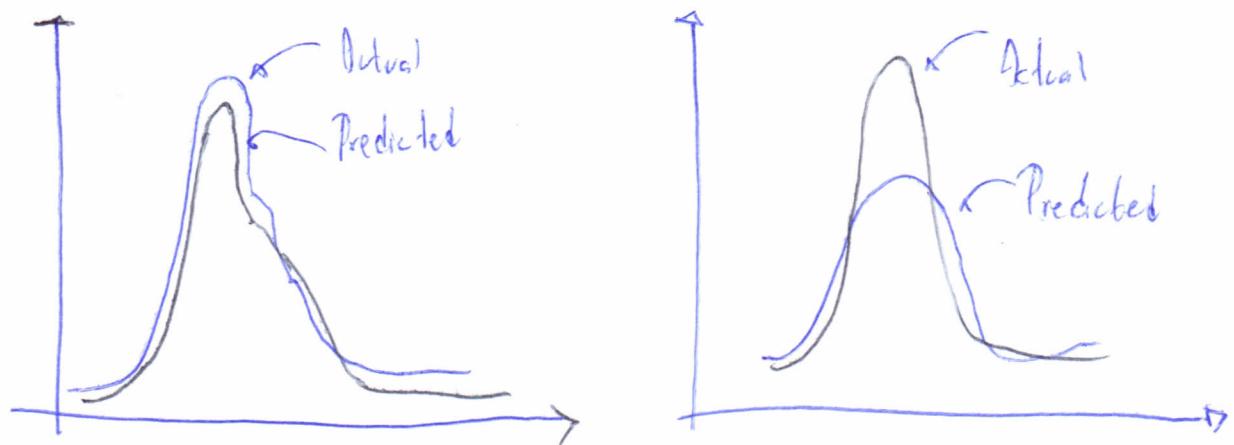
• Split dataset into:
 \hookrightarrow Training set (70%)
 \hookrightarrow Testing set (30%)

Data

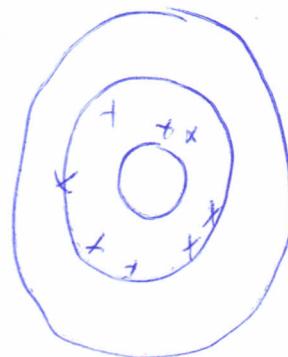
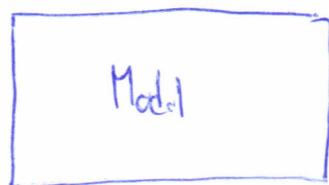
~~Mixed~~ Training / Testing Sets

Generalization Performance

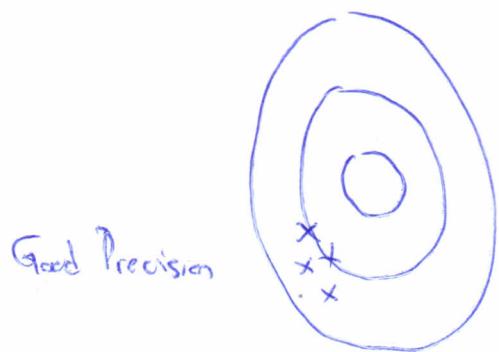
- Generalization error is measure of how well our data does at predicting previously unseen data.
- The error we obtain using our testing data is an approximation of this error.



Lots of Training Data



True Generalization Error



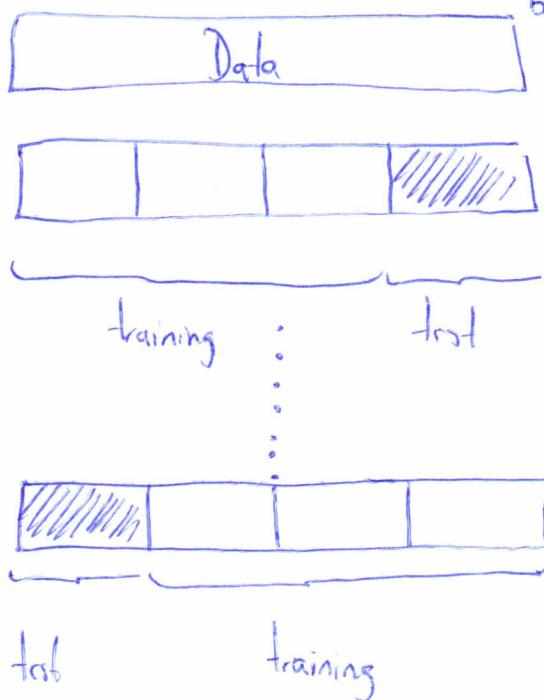
Good Precision

Cross Validation
to this problems

3

Cross-Validation

- Most common out-of-sample evaluation metrics
- More effective use of data (each observations is used for both training and testing)



Function `cross_val_score()`

```
from sklearn.model_selection import cross_val_score
```

```
scores = cross_val_score(lr, x_data, y_data, cv=3)
```

np.mean(scores)

predictor variable target variable

number of partitions

4

Function cross_val_score()



Model $\Rightarrow R^2 = \left(1 - \frac{\text{MSE of regression line}}{\text{MSE of } \bar{y}} \right)$



Function cross_val_predict()

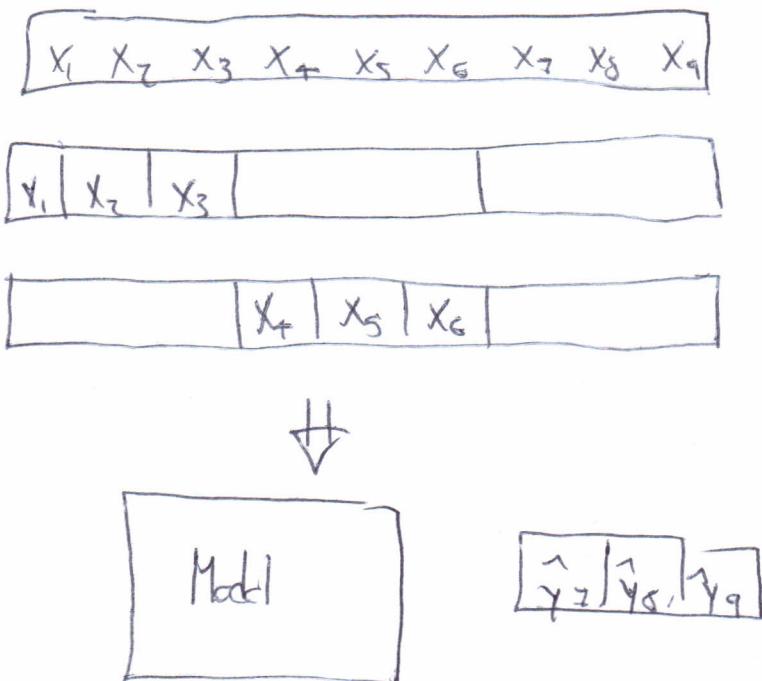
- It returns the prediction that was obtained for each element when it was in the test set
- Has a similar interface to cross_val_score()

```
from sklearn.model_selection import cross_val_predict
```

```
y_hat = cross_val_predict(lr2e, X_data, y_data, cv=3)
```

The Output is a prediction

Function cross_val_predict()

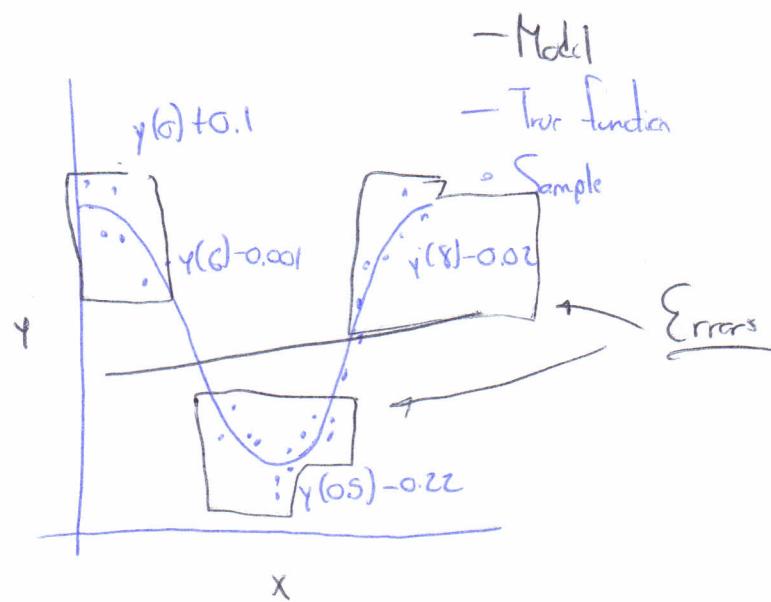


Overfitting, Underfitting and Model Selection

Model Selection

$$y(x) + \text{noise}$$

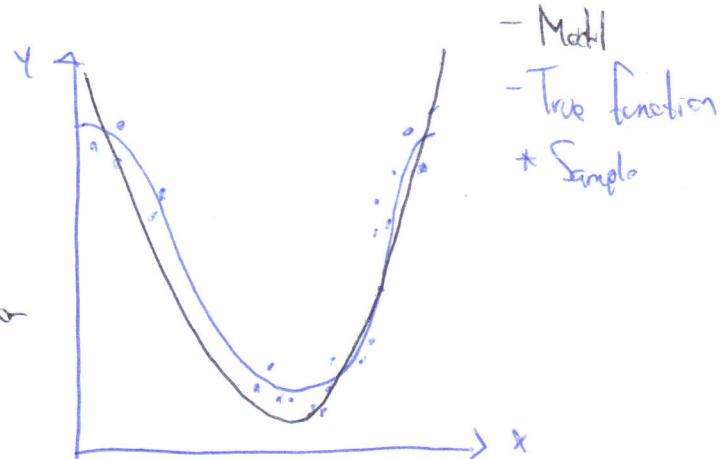
$$y = b_0 + b_1 x$$



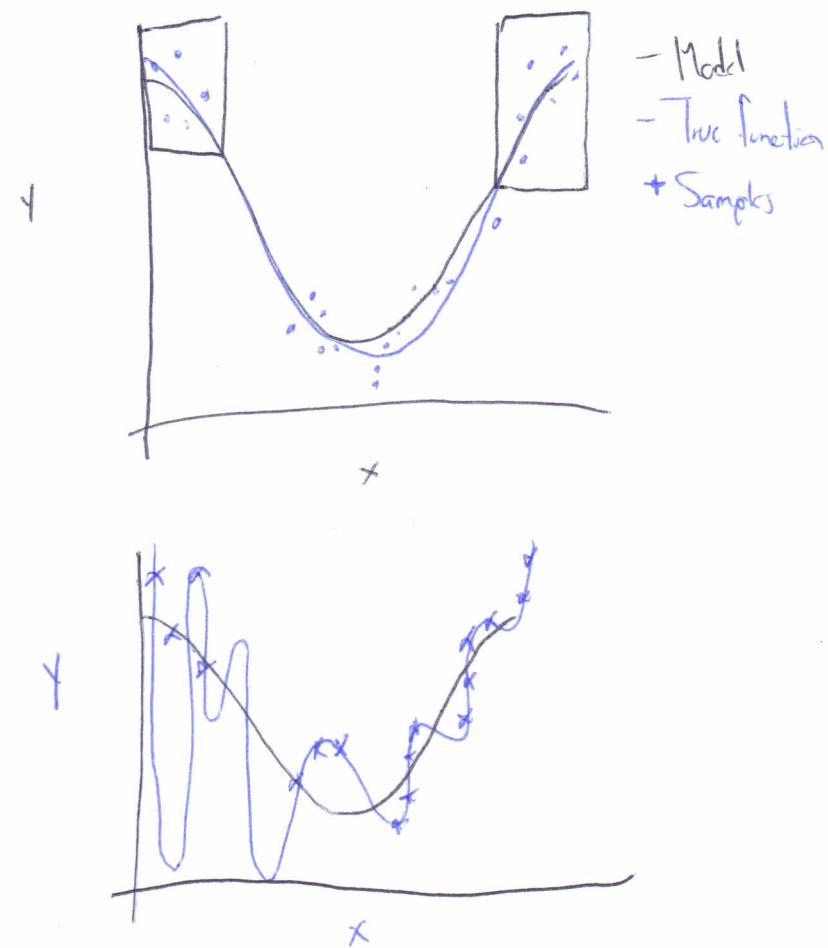
Underfitting \Rightarrow The model is too simple to fit the data.

$$y = b_0 + b_1 x + b_2 x^2$$

Increase the order \Rightarrow fits better
but the model is still not flexible
enough and exhibits underfitting



$$\hat{y} = b_0 + b_1 x + b_2 x^2 + b_3 x^3 + b_4 x^4 + b_5 x^5 + b_6 x^6 + b_7 x^7 + b_8 x^8$$

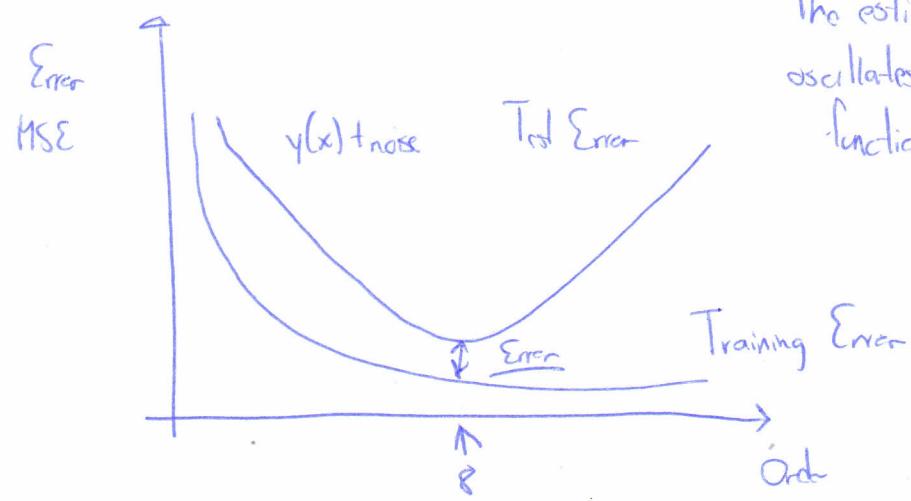


The estimated model is too flexible
and fits the noise rather than the function.

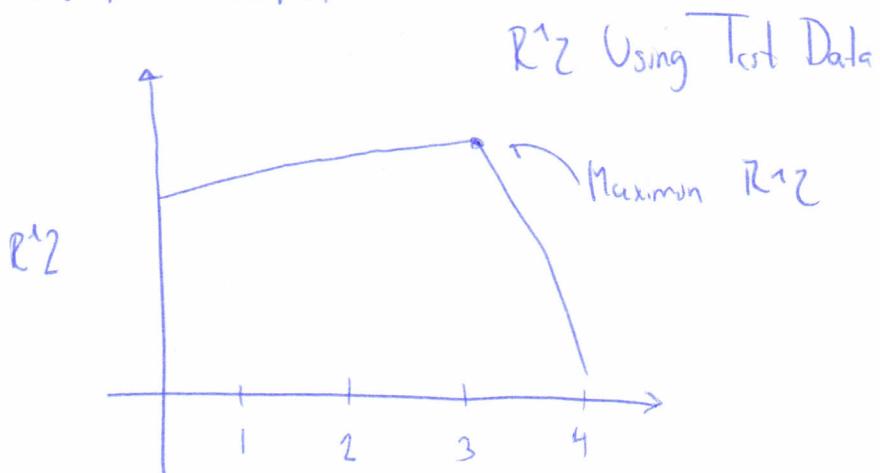
Over-fitting

\Downarrow
The estimated function
oscillates not tracking the
function.

Model Selection



Model Selection



```
R_squ_fct = []
order = [1, 2, 3, 4]
```

for n in order

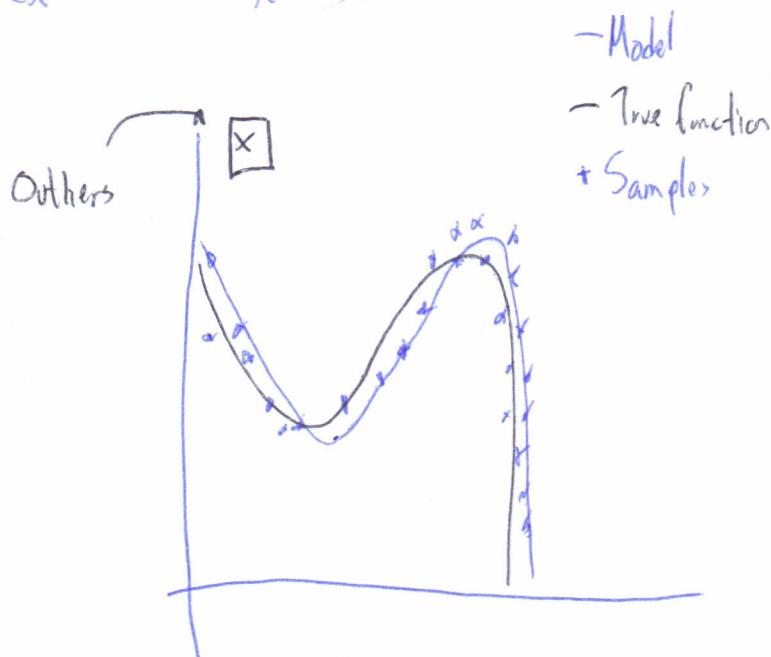
```
    pr = PolynomialFeatures(degree=n)
    x_train_pr = pr.fit_transform(x_train[['horsepower']])
    x_test_pr = pr.fit_transform(x_test[['horsepower']])
```

```
    lr.fit(x_train_pr, y_train)
```

```
Rsqu_fct.append((lr.score(x_test_pr, y_test)))
```

Ridge Regression

$$y = 1 + 2x - 3x^2 - 4x^3 + x^4$$



- Model

- True function

+ Samples

$$\hat{y} = 1 + 2x - 3x^2 - 2x^3 - 12x^4 - 40x^5 + 80x^6 + 71x^7 - 141x^8 - 38x^9 + 75x^{10}$$

Coefficients with a very large magnitude

Ridge Regression controls the magnitude of these polynomial coefficients.

<u>Alpha</u>	Parameter we select before fitting or training the model.
0	
0.001	
0.001	
1	
10	

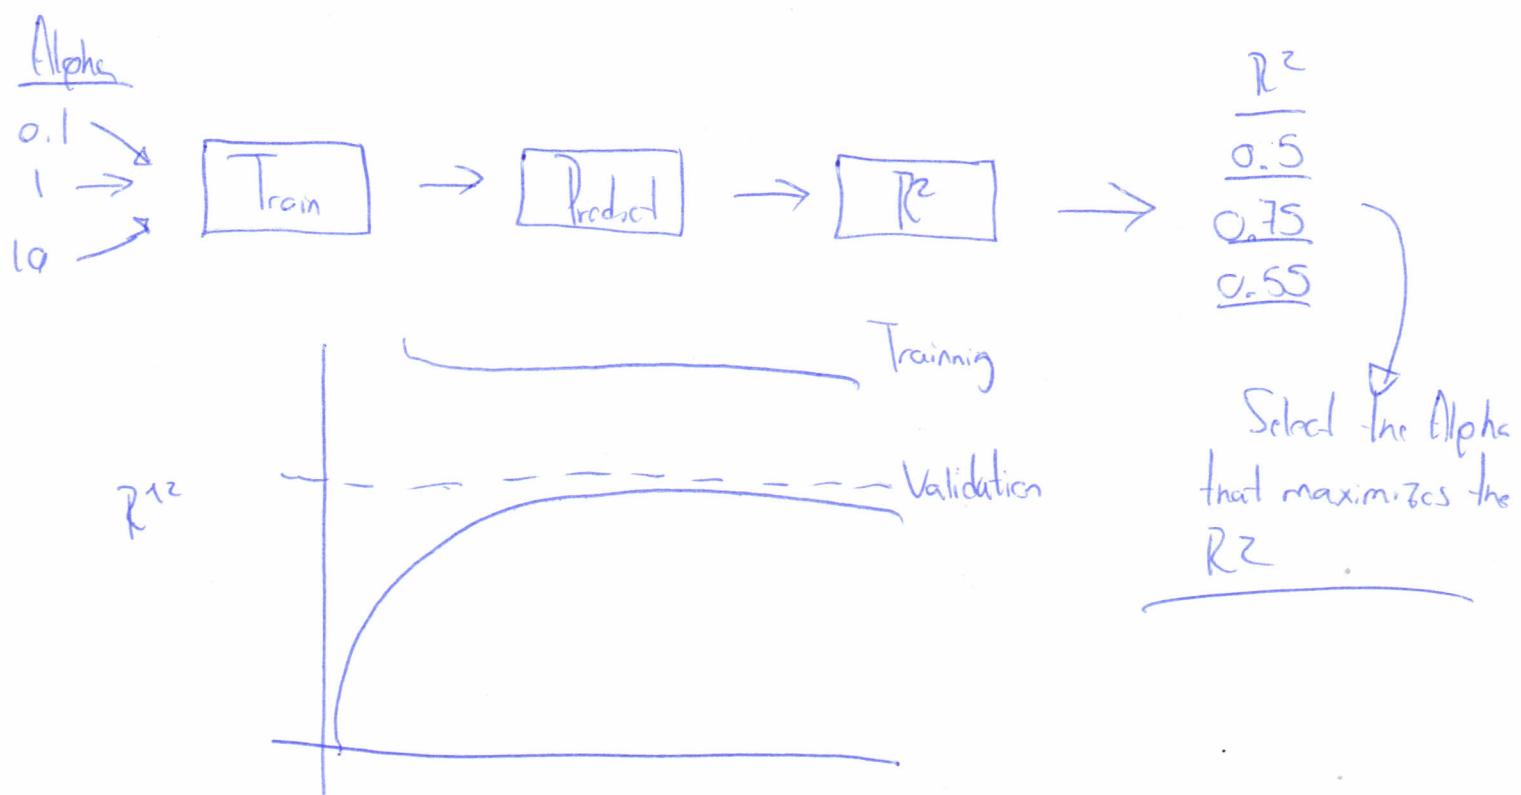
Ridge Regression

from `sklearn.linear_model import Ridge`

`RidgeModel = Ridge(alpha=0.1)`

`RidgeModel.fit(X, y)`

`Yhat = RidgeModel.predict(X)`



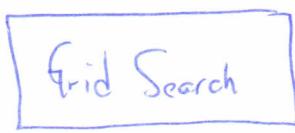
Overttting \Rightarrow Worse if we have a lots of features

Grid Search

Hyperparameters

{ θ ϕ ψ }

{ θ ϕ \emptyset }



Model 1 \rightarrow Error 1

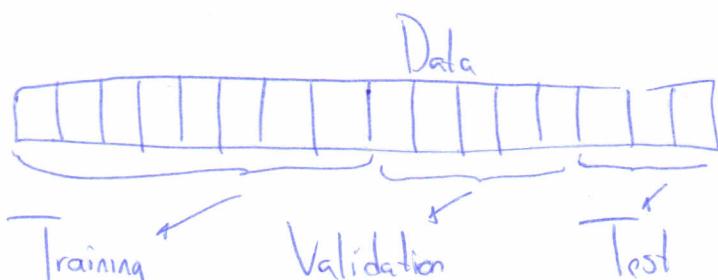
Model 2 \rightarrow Error 2

Model 3 \rightarrow Error 3

- In the last section, the term alpha in Ridge Regression is called a hyperparameter

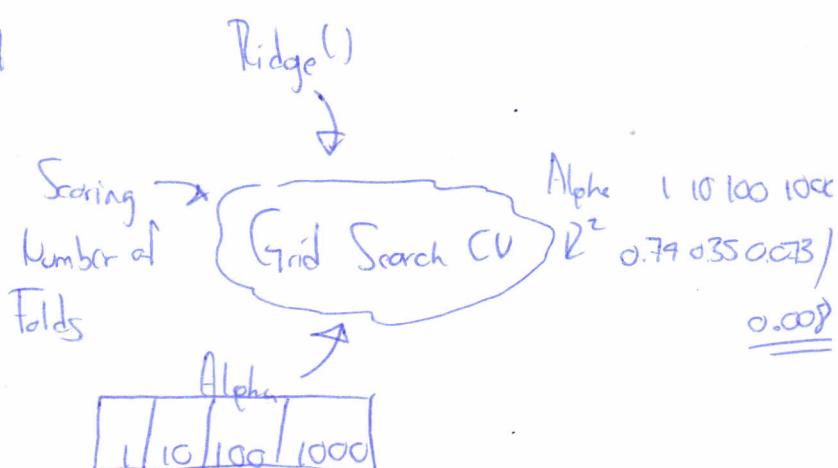
* Scikit-learn has a means of automatically iterating over these hyperparameters using cross-validation called Grid-Search

- Select the hyperparameters that minimize the error.



Hyperparameters

parameters = [{'alpha': [1, 10, 100, 1000]}]



Grid Search

```
from sklearn.linear_model import Ridge  
from sklearn.model_selection import GridSearchCV
```

```
parameters1 = [ {'alpha': [0.001, 0.1, 1, 10, 100, 1000, 10000, 100000]} ]
```

```
RR = Ridge()
```

```
Grid1 = GridSearchCV(RR, parameters1, cv=4)
```

```
Grid1.fit(x_data[['horsepower', 'curb-weight', 'engine-size', 'highway-mpg']], y_data)
```

```
Grid1.best_estimator_
```

```
scores = Grid1.cv_results_
```

```
scores['mean_test_scores']
```

```
parameters = [ {'alpha': [1, 10, 100, 1000], 'normalize': [True, False]} ]
```

Alpha	1	/	10	/	100	/	1000	/
Normalize	True		True		True		True	
False			False		False		False	

```
Ridge()
```

Ridge ()
↓



Alpha	1	10	100	1000
True	0.69	0.32	0.17	0.17
False	0.67	0.66	0.66	0.64

```
for param, mean_val, mean_test in zip(scores[params], scores['mean-test-score'], scores['Mean-train-scores'])
```

```
print(param, "R^2 on test data:", mean_val, "R^2 on train data:", mean_test)
```